

Baseball Case Study

Problem Statement:

This dataset utilizes data from 2014 Major League Baseball seasons in order to develop an algorithm that predicts the number of wins for a given team in the 2015 season based on several different indicators of success. There are 16 different features that will be used as the inputs to the machine learning and the output will be a value that represents the number of wins.

-- **Input features:** Runs, At Bats, Hits, Doubles, Triples, Homeruns, Walks, Strikeouts, Stolen Bases, *Runs* Allowed, Earned Runs, Earned Run Average (ERA), Shutouts, Saves, Complete Games and Errors

-- **Output:** Number of predicted wins (W)

Importing all the required libraries:

```
In [1]: import pandas as pd
import numpy as np
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

Importing the data:

```
In [2]: df=pd.read_csv("baseball.csv")
df.head()
```

```
Out[2]:
```

	W	R	AB	H	2B	3B	HR	BB	SO	SB	RA	ER	ERA	CG	SHO	SV	E
0	95	724	5575	1497	300	42	139	383	973	104	641	601	3.73	2	8	56	88
1	83	696	5467	1349	277	44	156	439	1264	70	700	653	4.07	2	12	45	86
2	81	669	5439	1395	303	29	141	533	1157	86	640	584	3.67	11	10	38	79
3	76	622	5533	1381	280	27	136	404	1231	68	701	643	3.98	7	9	37	101
4	74	689	5605	1515	289	49	151	455	1259	83	803	746	4.64	7	12	35	86

Data Analysis:

Checking if there are any null values present in the data

Checking for null values

```
In [3]: df.isnull().sum()
```

```
Out[3]: W      0
        R      0
        AB     0
        H      0
        2B     0
        3B     0
        HR     0
        BB     0
        SO     0
        SB     0
        RA     0
        ER     0
        ERA    0
        CG     0
        SHO    0
        SV     0
        E      0
        dtype: int64
```

Seems there are no null points present in the data

To check the shape that is to find the number of columns and rows in the data

```
In [4]: df.columns
```

```
Out[4]: Index(['W', 'R', 'AB', 'H', '2B', '3B', 'HR', 'BB', 'SO', 'SB', 'RA', 'ER',
               'ERA', 'CG', 'SHO', 'SV', 'E'],
              dtype='object')
```

To know the data in detailed we use describe funtion

```
In [5]: df.shape
```

```
Out[5]: (30, 17)
```

```
In [6]: df.describe()
```

```
Out[6]:
```

	W	R	AB	H	2B	3B	HR	
count	30.000000	30.000000	30.000000	30.000000	30.000000	30.000000	30.000000	30.0
mean	80.966667	688.233333	5516.266667	1403.533333	274.733333	31.300000	183.633333	469.1
std	10.453455	58.761754	70.467372	57.140923	18.095405	10.452355	31.823309	57.0
min	63.000000	573.000000	5385.000000	1324.000000	236.000000	13.000000	100.000000	375.0
25%	74.000000	651.250000	5484.000000	1383.000000	262.250000	23.000000	140.250000	428.2
50%	81.000000	689.000000	5510.000000	1382.500000	275.500000	31.000000	158.500000	473.0
75%	87.750000	718.250000	5570.000000	1451.500000	288.750000	39.000000	177.000000	501.2
max	100.000000	891.000000	5649.000000	1515.000000	308.000000	49.000000	232.000000	570.0

To check how many unique values are present in each column

```
In [7]: df.nunique()
```

```
Out[7]: W      24  
        R      28  
        AB     29  
        H      29  
        2B     22  
        3B     23  
        HR     27  
        BB     29  
        SO     29  
        SB     27  
        RA     30  
        ER     30  
        ERA    30  
        CG      9  
        SHO    12  
        SV     20  
        E      21  
dtype: int64
```

To check the datatype of each column present in the data

```
In [8]: df.dtypes
```


```
Out[8]: W          int64  
        R          int64  
        AB         int64  
        H          int64  
        2B         int64  
        3B         int64  
        HR         int64  
        BB         int64  
        SO         int64  
        SB         int64  
        RA         int64  
        ER         int64  
        ERA        float64  
        CG         int64  
        SHO        int64  
        SV         int64  
        E          int64  
dtype: object
```

Below are the catplot graphs where each column is compared with the WIN column

```
In [9]: import seaborn as sns
```


```
In [10]: sns.catplot(x='R',y='W',data=df)
```

```
Out[10]: <seaborn.axisgrid.FacetGrid at 0x238ffd3a430>
```




```
In [11]: sns.catplot(x='AB',y='W',data=df)
```

```
Out[11]: <seaborn.axisgrid.FacetGrid at 0x238ffe79220>
```




```
In [12]: sns.catplot(x='H',y='W',data=df)
```

```
Out[12]: <seaborn.axisgrid.FacetGrid at 0x238c0038be0>
```




```
In [13]: sns.catplot(x='28',y='W',data=df)
```

```
Out[13]: <seaborn.axisgrid.FacetGrid at 0x238ffe03730>
```




```
In [14]: sns.catplot(x='3B',y='W',data=df)
```

```
Out[14]: <seaborn.axisgrid.FacetGrid at 0x238c0132ca0>
```




```
In [15]: sns.catplot(x='HR',y='W',data=df)
```

```
Out[15]: <seaborn.axisgrid.FacetGrid at 0x238c02ae730>
```




```
In [16]: sns.catplot(x='BB',y='W',data=df)
```

```
Out[16]: <seaborn.axisgrid.FacetGrid at 0x238c02f08e0>
```




```
In [17]: sns.catplot(x='S0',y='W',data=df)
```

```
Out[17]: <seaborn.axisgrid.FacetGrid at 0x238c02a2a90>
```





```
In [18]: sns.catplot(x='SB',y='W',data=df)
```

```
Out[18]: <seaborn.axisgrid.FacetGrid at 0x238c042d7f0>
```




```
In [19]: sns.catplot(x='RA',y='W',data=df)
```

```
Out[19]: <seaborn.axisgrid.FacetGrid at 0x238c05498e0>
```




```
In [20]: sns.catplot(x='ER',y='W',data=df)
```

```
Out[20]: <seaborn.axisgrid.FacetGrid at 0x238c1680a0>
```




```
In [21]: sns.catplot(x='ERA',y='W',data=df)
```

```
Out[21]: <seaborn.axisgrid.FacetGrid at 0x238c175bdf0>
```




```
In [22]: sns.catplot(x='CG',y='W',data=df)
```

```
Out[22]: <seaborn.axisgrid.FacetGrid at 0x238c183d3d0>
```




```
In [23]: sns.catplot(x='SHO',y='W',data=df)
```

```
Out[23]: <seaborn.axisgrid.FacetGrid at 0x238ffe8a250>
```




```
In [24]: sns.catplot(x='SV',y='W',data=df)
```

```
Out[24]: <seaborn.axisgrid.FacetGrid at 0x238c15e94c0>
```



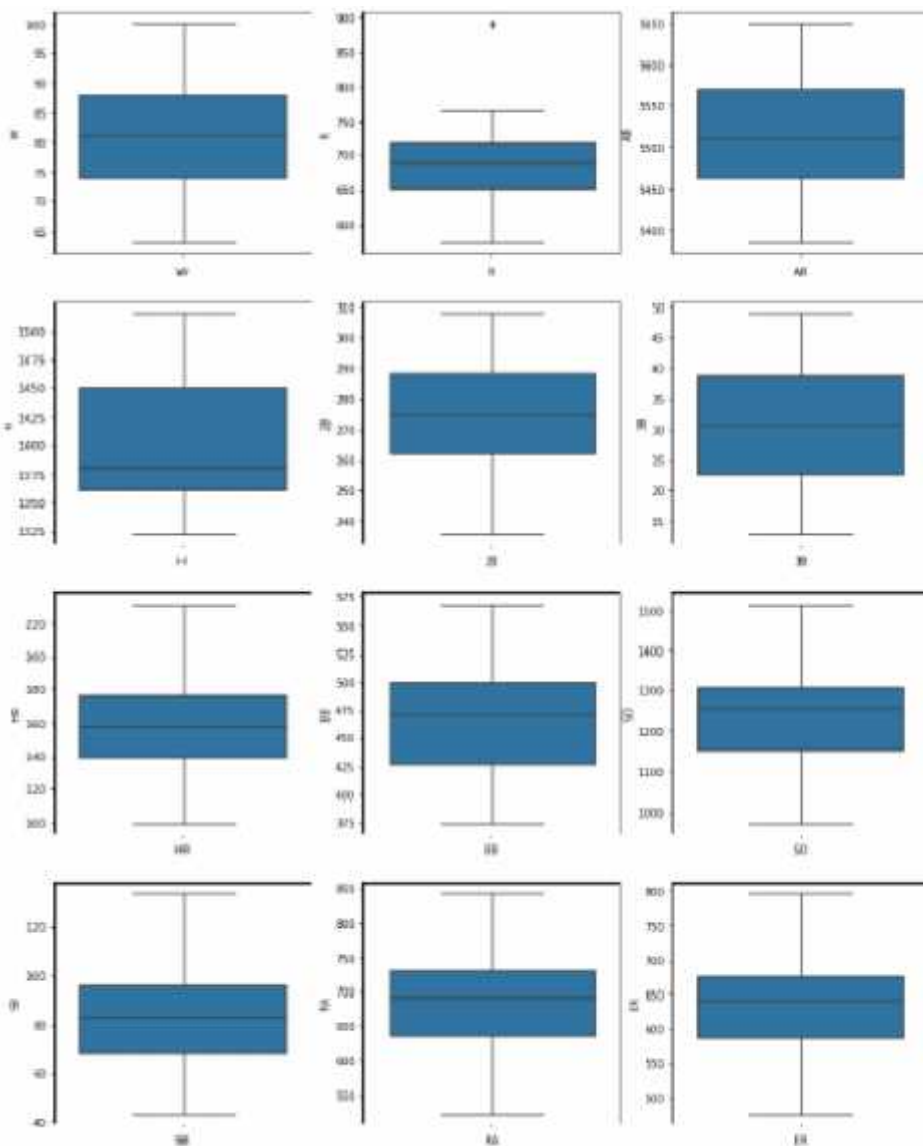
```
In [25]: sns.catplot(x='F',y='W',data=df)
```


```
Out[25]: <seaborn.axisgrid.FacetGrid at 0x238c19abc40>
```



Visualizing the outliers using the boxplot:

```
In [26]: #Visualize the outliers using boxplot
import matplotlib.pyplot as plt
plt.figure(figsize=(15,50))
graph=1
for column in df:
    if graph<=30:
        ax=plt.subplot(10,3,graph)
        sns.boxplot(df[column],orient='v')
        plt.xlabel(column,fontsize=10)
    graph+=1
plt.show()
```





Removing the outliers by using z-score

```
In [27]: #for removing outliers implementing zscore
from scipy.stats import zscore
z_score=zscore(df[['R', 'ERA', 'SHO', 'SV', 'E']])
abs_z_score=np.abs(z_score)

filtering_entry=(abs_z_score<3).all(axis=1)

df=df[filtering_entry]
```

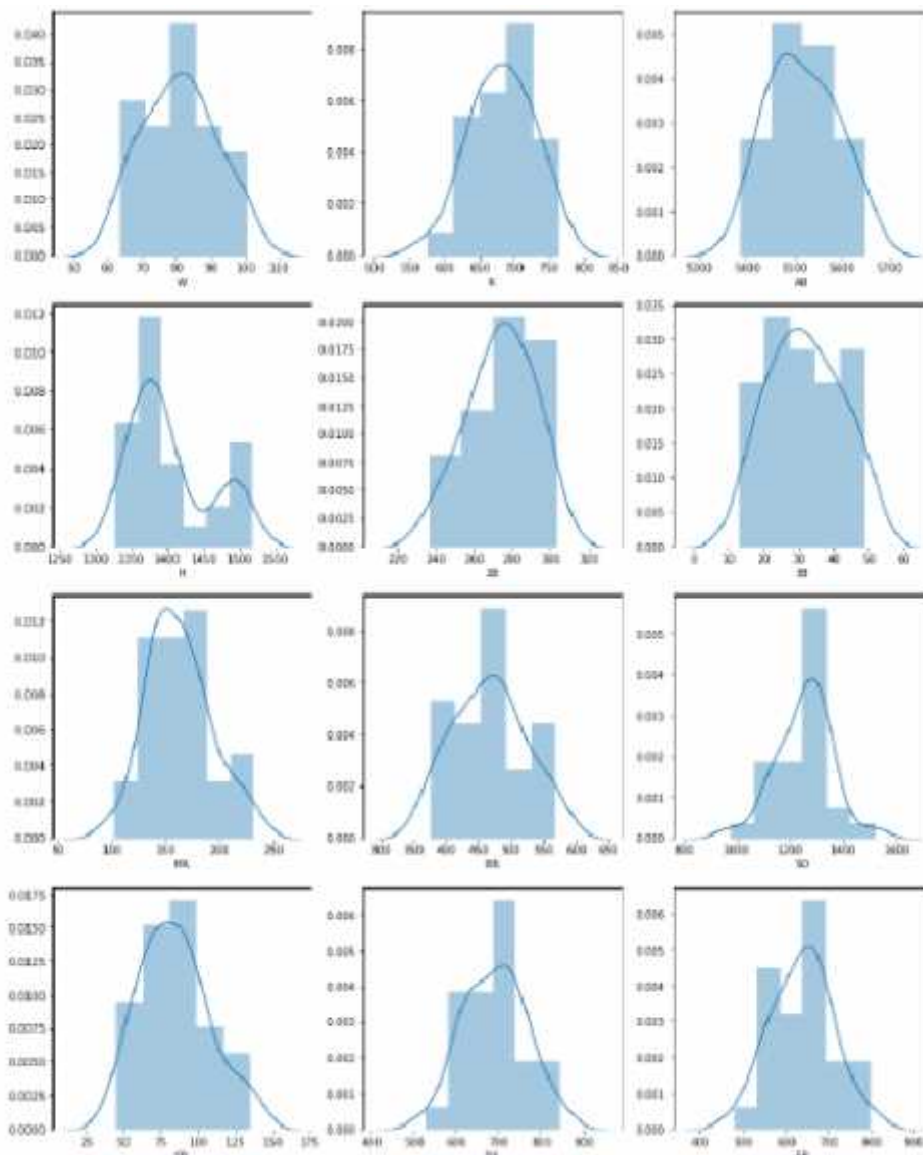
After removing the outliers the data ends up with 29 rows and 17 columns


```
In [31]: df.shape
```

```
Out[31]: (29, 17)
```

After removing the outliers, plotted the dist plot to check whether all columns have normal distribution

```
In [32]: import matplotlib.pyplot as plt
plt.figure(figsize=(15,50))
graph=1
for column in df:
    if graph<=30:
        ax=plt.subplot(10,3,graph)
        sns.distplot(df[column])
        plt.xlabel(column,fontsize=10)
        graph+=1
plt.show()
```





Now plotting the correlation graph to check the correlation of the data

```
In [33]: df.corr()
```

Out[33]:

	W	R	AB	H	2B	3B	HR	BB	SO
W	1.000000	0.390451	-0.085780	-0.018360	0.384886	-0.206737	0.245697	0.447513	0.156489
R	0.390451	1.000000	0.438022	0.433525	0.469293	0.134204	0.586894	0.258450	0.081158
AB	-0.085780	0.438022	1.000000	0.769159	0.490752	0.445604	-0.064653	-0.137850	-0.111243
H	-0.018360	0.433525	0.769159	1.000000	0.528016	0.582024	-0.218711	-0.222271	-0.371861
2B	0.384886	0.469293	0.490752	0.528016	1.000000	0.342419	-0.098695	0.211243	-0.096772
3B	-0.206737	0.134204	0.445604	0.582024	0.342419	1.000000	-0.369299	-0.404652	-0.196586
HR	0.245697	0.586894	-0.064653	-0.218711	-0.098695	-0.369299	1.000000	0.336614	0.479914
BB	0.447513	0.258450	-0.137850	-0.222271	0.211243	-0.404652	0.336614	1.000000	0.315566
SO	0.156489	0.081158	-0.111243	-0.371861	-0.096772	-0.196586	0.479914	0.315566	1.000000
SB	-0.169503	0.075323	0.373674	0.417877	0.194308	0.483818	-0.166072	-0.117622	0.038146
RA	-0.823176	-0.013858	0.315499	0.244606	-0.215196	0.312750	-0.092586	-0.425381	-0.140503
ER	-0.815308	0.007727	0.309146	0.280571	-0.224993	0.333731	-0.062094	-0.455832	-0.177604
ERA	-0.828952	-0.009122	0.254872	0.256458	-0.248212	0.325853	-0.070756	-0.465794	-0.195305
CG	0.029594	0.101438	-0.078511	0.092577	0.244856	-0.003733	0.065978	0.417437	-0.053000
SHO	0.497526	-0.065108	-0.196872	-0.135116	0.084060	-0.058896	0.005546	0.473922	0.231200
SV	0.749290	0.061381	-0.113342	-0.079814	0.269999	-0.210627	0.066984	0.187101	0.091300
E	-0.072858	0.043123	0.316297	-0.011945	0.145032	0.108610	-0.189790	-0.050114	0.142700

Here x is assigned to the input features of the model whereas y is assigned to the target.

```
In [35]: #Dividing data set into features and labels
y=df['W']
x=df.drop(columns=['W'])
```

ANOVA:

Analysis of variance (ANOVA) is an analysis tool used in statistics that splits an observed aggregate variability found inside a data set into two parts: systematic factors and random factors

```
In [36]: from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import f_classif
s=SelectKBest(f_classif,k=15)
s.fit(x,y)
anova=pd.DataFrame([s.scores_,s.pvalues_],columns=x.columns).T.sort_values(by=0)
```

```
In [37]: #ANOVA is used to determine the influence that independent variables have on the  
         andva
```

Out[37]:

	0	1
CG	0.361597	0.962964
H	0.729450	0.730204
2B	0.799063	0.680704
3B	0.811129	0.672290
HR	0.818974	0.666351
BB	0.943327	0.584607
SHO	1.253358	0.418187
SO	1.519009	0.316179
AB	1.622586	0.284961
ER	1.636442	0.201042
ERA	1.732208	0.255652
R	2.486509	0.130541
E	2.492758	0.129775
RA	2.524516	0.126473
SV	2.941436	0.091839
SB	3.283197	0.072181

DATA SCALING:

Data scaling is a method used to normalize the range of independent variables or features of data. In data processing, it is also known as data normalization and is generally performed during the data preprocessing step.

Data Scaling

```
In [38]: from sklearn.preprocessing import StandardScaler  
         scaler=StandardScaler()  
         x_scaled=scaler.fit_transform(x)
```

Building Machine Learning Models:

In the below piece of code, there is a function called classify where the data splitting is done , model fitting is done and accuracy of the model is calculated.

LogisticRegression , LinearRegression , DecisionTreeClassifier , RandomForestClassifier , ExtraTreeClassifier , XGBClassifier are implemented by calling the classify function. After fitting with everything mentioned above, The highest accuracy is shown by LinearRegression

```
In [40]: #Model training
from sklearn.model_selection import train_test_split, cross_val_score
def classify(model):
    x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=42)
    model.fit(x_train, y_train)
    print("Accuracy:", model.score(x_test, y_test))
```

```
In [42]: # Linear Regression model
from sklearn.linear_model import LinearRegression
model = LinearRegression()
classify(model)
```

Accuracy: 0.8906756651558634

The below graph is the predicted vs actual outcome graph

```
In [47]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=42)
```

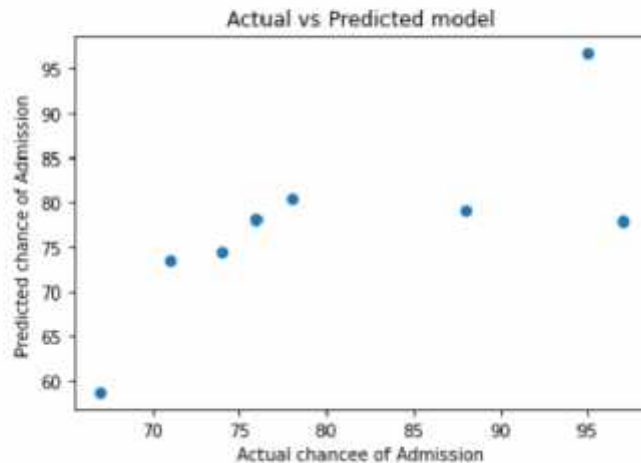
```
In [48]: model = LinearRegression()
model.fit(x_train, y_train)
```

```
Out[48]: LinearRegression()
```

```
In [49]: y_predicted = model.predict(x_test)
y_predicted
```

```
Out[49]: array([ 74.52025213,  77.84775591,  78.02499087,  58.71913825,  80.44146018,
  78.96373701,  73.3939036 ,  96.67416429])
```

```
In [50]: plt.scatter(y_test,y_predicted)
plt.xlabel('Actual chancee of Admission')
plt.ylabel('Predicted chance of Admission')
plt.title('Actual vs Predicted model')
plt.show()
```



Cross Validation and GridSearchCV is done

Cross-validation is a technique in which we train our model using the subset of the data-set and then evaluate using the complementary subset of the data-set.

```
In [51]: lm = LinearRegression()
from sklearn.model_selection import RepeatedKFold
rkf = RepeatedKFold(n_splits=4,n_repeats=2,random_state=True)

from sklearn.model_selection import cross_val_score
scores = cross_val_score(lm,x,y,cv=rkf)
print(scores)
print("Average score %.2f" % scores.mean())

[ 0.39669142  0.51602666  0.74669849 -0.29314478 -0.38619091 -0.63323544
  0.34375239  0.84459241]
Average score 0.19
```

GridSearchCV is a useful tool to fine tune the parameters of your model.

```
In [52]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Ridge
ridge_params={'alpha':[1,2,3,4,5,6,7,8,9,10]}
xg_grid=GridSearchCV(Ridge(),ridge_params,cv=3)
xg_grid.fit(x_train,y_train)
print('Best Score:',xg_grid.best_score_)
print('Best Score:',xg_grid.best_params_)
print('Best Score:',xg_grid.best_estimator_)
```

```
Best Score: -1.0850097103295599
Best Score: {'alpha': 10}
Best Score: Ridge(alpha=10)
```

Below is the piece of code to save the best model as a pickle file

```
In [53]: import pickle
model=LinearRegression()
model.fit(x_train, y_train)
pickle.dump(model, open('baseball.pkl','wb'))
```

```
In [54]: # Loading model to compare the results
loaded = pickle.load(open('baseball.pkl','rb'))
```

```
In [55]: y_predicted = model.predict(x_test)
y_predicted
```

```
Out[55]: array([74.52025213, 77.84775591, 78.02499087, 58.71913825, 80.44146018,
78.96373701, 73.3939036 , 96.67416429])
```

Concluding Remarks:

As discussed in EDA and storytelling the randomness of the dataset should be reduced by reducing the number of variables. Which was done by Cross Validation and GridSearchCV.